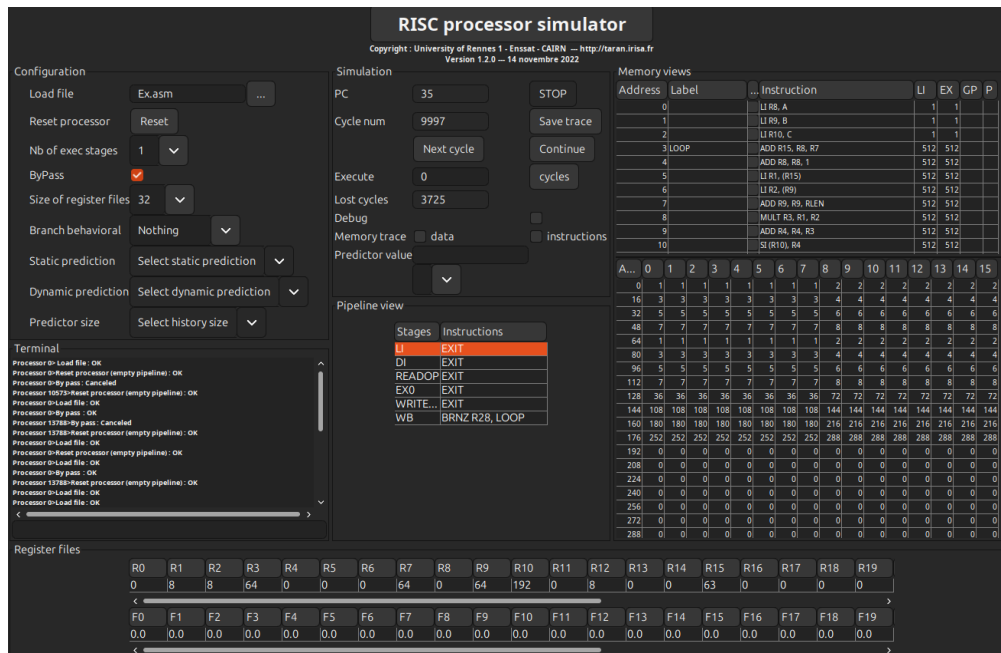


# Nguyen Tai Anh- 22BI13028- Practical 3- Advanced CA and x86 ISA

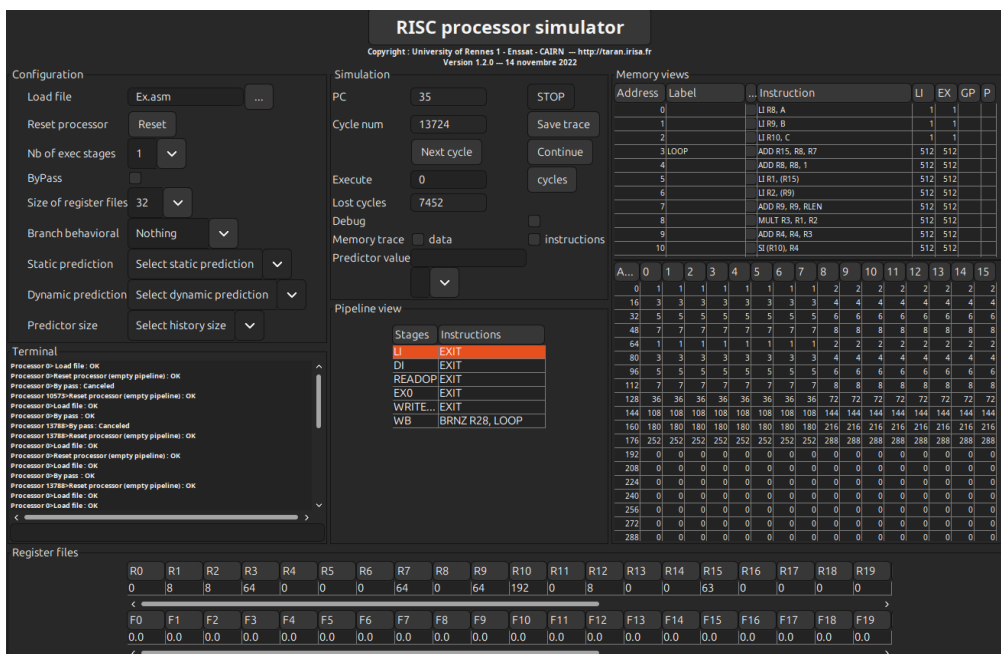
## Step 1: Analyze Bypass technique

- Execute code without and with bypass

With ByPass:



Without ByPass:



What is ByPass technique:

+ Bypassing, also known as "data forwarding," is a technique used in modern CPU architectures, including x86, to improve the efficiency of pipelined processors. It helps to reduce the delay caused by data hazards, specifically read-after-write (RAW) hazards, in the instruction pipeline.

+ Bypassing addresses this issue by allowing an instruction to use the result of a previous instruction as soon as it is computed, without waiting for it to be written back to the register file. This is achieved by forwarding the result directly from one stage of the pipeline to another. Typically, the result is forwarded from the output of the execution stage (EX) to the input of the execution stage of the dependent instruction.

| Stages   | Instructions     |
|----------|------------------|
| LI       | ADD R4, R4, R3   |
| DI       | MULT R3, R1, R2  |
| READOP   | ADD R9, R9, RLEN |
| EX0      | LI R2, (R9)      |
| WRITE... | LI R1, (R15)     |
| WB       | ADD R8, R8, 1    |

| Stages   | Instructions     |
|----------|------------------|
| LI       | MULT R3, R1, R2  |
| DI       | ADD R9, R9, RLEN |
| READOP   | LI R2, (R9)      |
| EX0      | LI R1, (R15)     |
| WRITE... | -                |
| WB       | ADD R8, R8, 1    |

As can be seen from the two images above, on the left, ByPass is enabled and on the right, ByPass is disabled, we can see the main difference between them. After calculating R15 from R8, the result of R15 immediately directs to Li R1, (R15) and executes the instruction when ByPass is on. But when ByPass is disabled, there will be a lost cycle between ADD R8, R8, 1 and Li R1, (R15), means that after calculating R15 from R8, the result of R15 must be write to the register first and then next direct to Li R1, (R15) later ( or Li R1, (R15) have to wait for the previous instruction to be written first).

### - Analyze the results (nb of cycles), comments, conclusion

With ByPass: There are 9997 cycle numbers, and 3725 of them are lost during the execution.

Without ByPass: There are 13724 cycle numbers, and 7452 of them are lost during the execution.

→ As can be seen, without ByPass, there are more cycles as well as lost cycles during execution.

→ With bypass, the pipeline can avoid the stalls forward the result of the first instruction to the second instruction (The result of the first instruction will directly transfer to the second instruction )→ no waiting for the result written to register file and read back in → instructions are executed more quickly and efficiently.

### - Modify the nb of pipeline execution stages, and analyze what happen



**Note:** I only show the ByPass enabled case when changing the nb of exec stages!!!

Conclusion: As can be seen from these pictures above, I set the number of execution stages from 1, 2, 4, 6, 10 → when the number of execution stages increases, the cycle numbers and lost cycles also increase. But it seems like the lost cycles increase faster and more compared to the cycle numbers.

| Number of<br>exec stage(s) | ByPass enabled |             | ByPass disabled |             |
|----------------------------|----------------|-------------|-----------------|-------------|
|                            | Cycle numbers  | Lost cycles | Cycle numbers   | Lost cycles |
| 1                          | 9997           | 3725        | 13742           | 7452        |
| 2                          | 13725          | 7452        | 17453           | 11180       |
| 4                          | 21183          | 14908       | 24911           | 18636       |
| 6                          | 28697          | 22420       | 32425           | 26148       |
| 10                         | 44061          | 37780       | 47789           | 41508       |

Explanation: Increasing the number of execution stages in a pipeline can improve instruction throughput by allowing for higher clock speeds and more parallel instruction processing (means that the processor will work more and reduce the execution time). However, this also introduces greater complexity in managing pipeline hazards, increases the branch misprediction penalty, and adds overhead for pipeline control. These factors contribute to a higher number of total cycles and lost cycles, as the processor works to handle the additional challenges introduced by the deeper pipeline.

## **Step 2: Analyze branch prediction techniques (static)**

**- Configure static prediction + Analyze the nb of cycles, comments, conclusion**

**+, “+/- not taken”:**

**RISC processor simulator**  
Copyright : University of Rennes 1 - Enssat - CAIRN — <http://taran.lriia.fr>  
Version 1.2.0 — 14 novembre 2022

**Configuration**

Load file:  ...

Reset processor:

Nb of exec stages:  ▼

ByPass: ☒

Size of register files:  ▼

Branch behavioral:  ▼

Static prediction:  ▼

Dynamic prediction:  ▼

Predictor size:  ▼

**Terminal**

```

Processor 0- Load file : OK
Processor 0-Reset processor (empty pipeline) : OK
Processor 0-Reset processor (empty pipeline) : OK
Processor 0-Load file : OK
Processor 0-By pass : Cancelled
Processor 0-By pass : OK
Processor 13728-Reset processor (empty pipeline) : OK
Processor 0-Load file : OK
Processor 0-Configuration of pipeline : OK
Processor 0-Reset processor (empty pipeline) : OK
Processor 0-Configuration of pipeline : OK
Processor 13728-Reset processor (empty pipeline) : OK
Processor 0-Reset processor (empty pipeline) : OK
Processor 0-Load file : OK
Processor 0-Configuration of pipeline : OK

```

**Simulation**

PC:

Cycle num:

Execute:

Lost cycles:

Debug: ☐

Memory trace: ☐ data ☐ instructions

Predictor value:  ▼

**Pipeline view**

| Stages   | Instructions   |
|----------|----------------|
| LI       | EXIT           |
| DI       | EXIT           |
| READOP   | EXIT           |
| EX0      | EXIT           |
| WRITE... | EXIT           |
| WB       | BRNZ R28, LOOP |

**Memory views**

| Address | Label | Instruction      | LI  | EX  | GP | P |
|---------|-------|------------------|-----|-----|----|---|
| 0       |       | LI R8, A         | 1   | 1   |    |   |
| 1       |       | LI R9, B         | 1   | 1   |    |   |
| 2       |       | LI R10, C        | 1   | 1   |    |   |
| 3       | LOOP  | ADD R15, R8, R7  | 512 | 512 |    |   |
| 4       |       | ADD R8, R8, 1    | 512 | 512 |    |   |
| 5       |       | LI R1, (R15)     | 512 | 512 |    |   |
| 6       |       | LI R2, (R9)      | 512 | 512 |    |   |
| 7       |       | ADD R9, R9, R1EN | 512 | 512 |    |   |
| 8       |       | MULT R3, R1, R2  | 512 | 512 |    |   |
| 9       |       | ADD R4, R4, R3   | 512 | 512 |    |   |
| 10      |       | IS (R10), R4     | 512 | 512 |    |   |

| A... | 0   | 1   | 2   | 3   | 4   | 5   | 6   | 7   | 8   | 9   | 10  | 11  | 12  | 13  | 14  | 15  |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0    | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 2   | 2   | 2   | 2   | 2   | 2   | 2   | 2   |
| 16   | 3   | 3   | 3   | 3   | 3   | 3   | 3   | 3   | 4   | 4   | 4   | 4   | 4   | 4   | 4   | 4   |
| 32   | 5   | 5   | 5   | 5   | 5   | 5   | 5   | 5   | 6   | 6   | 6   | 6   | 6   | 6   | 6   | 6   |
| 48   | 7   | 7   | 7   | 7   | 7   | 7   | 7   | 7   | 8   | 8   | 8   | 8   | 8   | 8   | 8   | 8   |
| 64   | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 2   | 2   | 2   | 2   | 2   | 2   | 2   | 2   |
| 80   | 3   | 3   | 3   | 3   | 3   | 3   | 3   | 3   | 4   | 4   | 4   | 4   | 4   | 4   | 4   | 4   |
| 96   | 5   | 5   | 5   | 5   | 5   | 5   | 5   | 5   | 6   | 6   | 6   | 6   | 6   | 6   | 6   | 6   |
| 112  | 7   | 7   | 7   | 7   | 7   | 7   | 7   | 7   | 8   | 8   | 8   | 8   | 8   | 8   | 8   | 8   |
| 128  | 36  | 36  | 36  | 36  | 36  | 36  | 36  | 36  | 72  | 72  | 72  | 72  | 72  | 72  | 72  | 72  |
| 144  | 108 | 108 | 108 | 108 | 108 | 108 | 108 | 108 | 144 | 144 | 144 | 144 | 144 | 144 | 144 | 144 |
| 160  | 180 | 180 | 180 | 180 | 180 | 180 | 180 | 180 | 216 | 216 | 216 | 216 | 216 | 216 | 216 | 216 |
| 176  | 252 | 252 | 252 | 252 | 252 | 252 | 252 | 252 | 288 | 288 | 288 | 288 | 288 | 288 | 288 | 288 |
| 192  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 208  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 224  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 240  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 256  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 272  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 288  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |

**Register files**

| R0 | R1 | R2 | R3 | R4 | R5 | R6 | R7 | R8 | R9 | R10 | R11 | R12 | R13 | R14 | R15 | R16 | R17 | R18 | R19 |
|----|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0  | 8  | 8  | 64 | 0  | 0  | 0  | 64 | 0  | 64 | 192 | 0   | 8   | 0   | 0   | 63  | 0   | 0   | 0   | 0   |

| F0  | F1  | F2  | F3  | F4  | F5  | F6  | F7  | F8  | F9  | F10 | F11 | F12 | F13 | F14 | F15 | F16 | F17 | F18 | F19 |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |

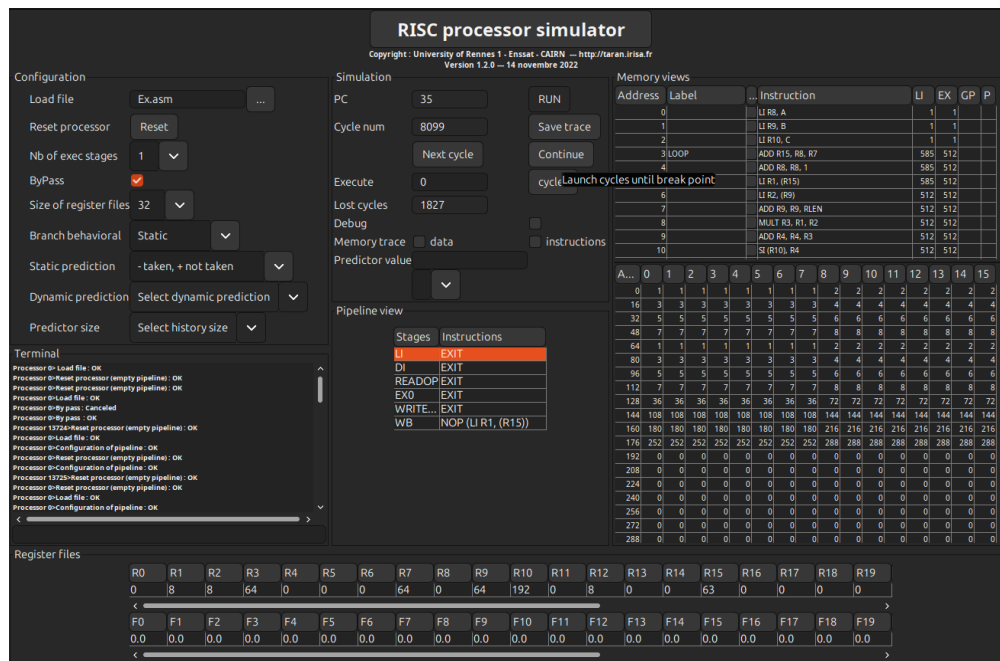
The “+/- not taken” means that for backward and forward conditional branch instructions, the processor will try to NOT jump. As can be seen from the image, when I enable Static branch behavior and set the Static prediction to “+/- not taken”, the Cycle num and Lost cycles slightly decreases to 9413 and 3141

**+, “+/- taken”:**



“- not taken, + taken” means that for backward conditional branch instructions, the processor will try to NOT jump and for forward conditional branch instructions, the processor will try to jump. As can be seen from the image, the result after execution is exactly the same as when executed with prediction “+/- not taken”. There are 9413 cycle num and 3141 lost cycles.

**+, “- taken, + not taken”:**



“- taken, + not taken” means that for backward conditional branch instructions, the processor will try to jump and for forward conditional branch instructions, the processor will try to NOT jump. As can be seen from the image, the result is the same as the situation when “+/- taken” with 8099 cycle num and 1827 lost cycles.

- The good configuration among these 4 predictions that can be seen is “+/- not taken” and “- taken, + not taken” as they are the clearest and have the least number of cycles as well as lost cycles. Reason why they have the same result is because in my code, there are only backward (or “-”) branch instructions.

### - Analyze cycle by cycle (- taken, + not taken)

In all my loops, there are only backward branch instructions and no forward branch instructions.

For loop1, loop2 and loop3:

```

Loop
    add R15, R8, R7
    add R8, R8, 1
    li R1, (R15)

    li R2, (R9)
    add R9, R9, RLEN

    mult R3, R1, R2
    add R4, R4, R3
    si (R10), R4

    sub R5, R5, 1
    add R6, R5, RLEN
    brnz R6, Loop

    add R10, R10, 1

```

| Stages   | Instructions        |
|----------|---------------------|
| LI       | ADD R15, R8, R7     |
| DI       | NOP (LI R5, 0)      |
| READOP   | NOP (LI R4, 0)      |
| EX0      | NOP (ADD R10, R1... |
| WRITE... | BRNZ R6, LOOP       |
| WB       | -                   |

**Not using Static  
branch prediction**

| Stages   | Instructions     |
|----------|------------------|
| LI       | LI R2, (R9)      |
| DI       | LI R1, (R15)     |
| READOP   | ADD R8, R8, 1    |
| EX0      | ADD R15, R8, R7  |
| WRITE... | BRNZ R6, LOOP    |
| WB       | ADD R6, R5, RLEN |

**Using Static  
branch prediction**

When configuring Static branch behavior and enabling Static mode to “- taken, + not taken”, the backward branch instruction will be predicted and remove all the NOP instructions that come after these backward branches while forward branch instructions will not be predicted. In this case, my code only has backward branch instructions.

As can be seen, in loop1, when Static branch prediction “- taken, + not taken” is disabled, after the “BRNZ R6, LOOP” instruction is written, there are 3 NOP instructions from loop2 are executed before instructions in loop1 are executed again.

But when Static branch prediction “- taken, + not taken” is enabled, after the instruction “BRNZ R6, LOOP” is written, instruction from loop1 is executed again immediately rather than 3 NOP instructions from loop2. → Reduce the cycle numbers and lost cycles during execution of the whole program.

For loop2 and loop3, there are also two BRNZ instructions → The result will be the same as in loop1. These branches will also skip the NOP instructions from their next loop and immediately come back to instructions in their loop.



```
.Loop2
    li R4, 0
    li R5, 0
    li R8, 0

    sub R9, R9, 63

    add R11, R11, 1
    sub R29, R11, RLEN
    brnz R29, Loop

    add R7, R7, RLEN
```

```

Loop3
    li R11, 0
    li R9, SIZE

    add R12, R12, 1
    sub R28, R12, RLEN

    brnz R28, Loop

```

### Step 3: Configure dynamic branch prediction, 1 bit (explain how you configure)

- **Configure dynamic branch prediction, 1-bit (explain how you configure)**

I change the Branch behavior to Dynamic and choose Dynamic prediction to 1-bit.

# RISC processor simulator

Copyright : University of Rennes 1 - Enssat - CARN — <http://taran.irisa.fr>  
Version 1.2.0 — 14 novembre 2022

### Configuration

Load file Ex.asm ...

Reset processor Reset

Nb of exec stages 1 ▼

ByPass ☒

Size of register files 32 ▼

Branch behavioral Dynamic ▼

Static prediction - taken, + not taken ▼

Dynamic prediction 1 bit ▼

Predictor size Select history size ▼

### Terminal

Processor 0-Branch delay configuration : OK

Processor 0-Load file : OK

Processor 0-Reset processor (empty pipeline) : OK

Processor 0-Run processor (empty pipeline) : OK

Processor 0-Load file : OK

Processor 0-Run processor (empty pipeline) : OK

Processor 0-Load file : OK

Processor 0-Reset processor (empty pipeline) : OK

Processor 0-Load file : OK

Processor 0-Dynamic branch configuration : OK

Processor 0-Static branch configuration : OK

Processor 0-Static branch configuration of the number of bits for dynamic prediction : OK

Processor 0-Static branch configuration of the number of bits for dynamic prediction : OK

Processor 0-Load file : OK

Processor 0-Run processor (empty pipeline) : OK

Processor 0-Load file : OK

### Simulation

PC 35 STOP

Cycle num 8309 Save trace

Next cycle Continue

Execute 0 cycles

Last cycles 2037

Debug ☐

Memory trace ☐ data ☐ instructions

Predictor value ▼

### Pipeline view

| Stages | Instructions   |
|--------|----------------|
| IF     | EXIT           |
| DI     | EXIT           |
| READOP | EXIT           |
| EXOP   | EXIT           |
| WRITE  | EXIT           |
| WB     | NOP (LR1, R15) |

### Memory views

| Address | Label     | Instruction      | U | EX | GP  | P   |
|---------|-----------|------------------|---|----|-----|-----|
| 0       | LR R8, A  |                  | 1 | 1  |     |     |
| 1       | LR R8, B  |                  | 1 | 1  |     |     |
| 2       | LR R10, C |                  | 1 | 1  |     |     |
| 3       | LOOP      | ADD R15, R8, R7  |   |    | 585 | 512 |
| 4       |           | ADD R8, R8, 1    |   |    | 585 | 512 |
| 5       |           | LR R1, (R15)     |   |    | 585 | 512 |
| 6       |           | LR R2, (R9)      |   |    | 512 | 512 |
| 7       |           | ADD R9, R9, R1EN |   |    | 512 | 512 |
| 8       |           | MULT R3, R1, R2  |   |    | 512 | 512 |
| 9       |           | ADD R4, R4, R3   |   |    | 512 | 512 |
| 10      |           | ST (R10), R4     |   |    | 512 | 512 |

| A... | 0   | 1   | 2   | 3   | 4   | 5   | 6   | 7   | 8   | 9   | 10  | 11  | 12  | 13  | 14  | 15  |
|------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0    | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 2   | 2   | 2   | 2   | 2   | 2   | 2   | 2   |
| 16   | 3   | 3   | 3   | 3   | 3   | 3   | 3   | 4   | 4   | 4   | 4   | 4   | 4   | 4   | 4   | 4   |
| 21   | 5   | 5   | 5   | 5   | 5   | 5   | 5   | 6   | 6   | 6   | 6   | 6   | 6   | 6   | 6   | 6   |
| 48   | 7   | 7   | 7   | 7   | 7   | 7   | 7   | 8   | 8   | 8   | 8   | 8   | 8   | 8   | 8   | 8   |
| 60   | 1   | 1   | 1   | 1   | 1   | 1   | 1   | 2   | 2   | 2   | 2   | 2   | 2   | 2   | 2   | 2   |
| 84   | 3   | 3   | 3   | 3   | 3   | 3   | 3   | 4   | 4   | 4   | 4   | 4   | 4   | 4   | 4   | 4   |
| 96   | 5   | 5   | 5   | 5   | 5   | 5   | 5   | 6   | 6   | 6   | 6   | 6   | 6   | 6   | 6   | 6   |
| 112  | 7   | 7   | 7   | 7   | 7   | 7   | 7   | 8   | 8   | 8   | 8   | 8   | 8   | 8   | 8   | 8   |
| 128  | 36  | 36  | 36  | 36  | 36  | 36  | 36  | 72  | 72  | 72  | 72  | 72  | 72  | 72  | 72  | 72  |
| 144  | 108 | 108 | 108 | 108 | 108 | 108 | 108 | 144 | 144 | 144 | 144 | 144 | 144 | 144 | 144 | 144 |
| 160  | 180 | 180 | 180 | 180 | 180 | 180 | 180 | 216 | 216 | 216 | 216 | 216 | 216 | 216 | 216 | 216 |
| 176  | 252 | 252 | 252 | 252 | 252 | 252 | 252 | 288 | 288 | 288 | 288 | 288 | 288 | 288 | 288 | 288 |
| 192  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 208  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 224  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 240  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 256  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 272  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |
| 288  | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   | 0   |

### Register files

### - Analyze the nb of cycles, comments, conclusion

As can be seen from the image, when I enable Dynamic branch behavior technique, the cycle number decreases to 8309 compared with more than 9900 originally. The lost cycles decrease significantly to 2037.

Explanation: Dynamic branch prediction is a technique used in modern processors to guess the outcome of a branch instruction (whether it will be taken or not) at runtime, rather than relying on a fixed strategy (static prediction). This technique uses historical information to make more accurate predictions, improving instruction pipeline efficiency and overall processor performance. That is why the lost cycles decrease so much as the cycle number decreases slightly.

### - Analyze cycle by cycle

In all my loops, there are only backward branch instructions and no forward branch instructions.

For loop1, loop2 and loop3:

```
.Loop
    add R15, R8, R7
    add R8, R8, 1
    li R1, (R15)

    li R2, (R9)
    add R9, R9, RLEN

    mult R3, R1, R2
    add R4, R4, R3
    si (R10), R4

    sub R5, R5, 1
    add R6, R5, RLEN
    brnz R6, Loop

    add R10, R10, 1
```

| Stages   | Instructions        |
|----------|---------------------|
| LI       | ADD R15, R8, R7     |
| DI       | NOP (LI R5, 0)      |
| READOP   | NOP (LI R4, 0)      |
| EX0      | NOP (ADD R10, R1... |
| WRITE... | BRNZ R6, LOOP       |
| WB       | -                   |

**Not using Dynamic  
branch prediction:  
1-bit mode**

| Stages   | Instructions     |
|----------|------------------|
| LI       | LI R2, (R9)      |
| DI       | LI R1, (R15)     |
| READOP   | ADD R8, R8, 1    |
| EX0      | ADD R15, R8, R7  |
| WRITE... | BRNZ R6, LOOP    |
| WB       | ADD R6, R5, RLEN |

**Using Dynamic  
branch prediction:  
1-bit mode**

A 1-bit branch predictor utilizes a single bit to remember the outcome of the last branch instruction. This bit, often referred to as a prediction bit, can be in one of two states: "taken" or "not taken." In the given scenario, the first bit would be set to 1 because when the instruction "BRNZ R6, LOOP" is executed, there are no intervening instructions such as "NOP (ADD R10, R10, 1)" or "NOP (LI R4, 0)" following it before reaching instructions within the Loop1 ("ADD R15, R8, R7"). Consequently, based on

this history, the predictor would predict that the branch will be taken in subsequent executions. During program execution, this prediction bit would be updated dynamically based on the actual outcomes of branch instructions, thereby refining the predictor's accuracy over time.

| Memory views |       |     |                  |    |    |    |   |
|--------------|-------|-----|------------------|----|----|----|---|
| Address      | Label | ... | Instruction      | LI | EX | GP | P |
| 9            |       |     | ADD R4, R4, R3   | 4  | 4  |    |   |
| 10           |       |     | SI (R10), R4     | 4  | 3  |    |   |
| 11           |       |     | SUB R5, R5, 1    | 4  | 3  |    |   |
| 12           |       |     | ADD R6, R5, RLEN | 4  | 3  |    |   |
| 13           |       |     | BRNZ R6, LOOP    | 3  | 3  | 3  | T |
| 14           |       |     | ADD R10, R10, 1  | 0  | 0  |    |   |
| 15           |       |     | BR LOOP2         | 0  | 0  |    |   |
| 16           | LOOP2 |     | LI R4, 0         | 0  | 0  |    |   |
| 17           |       |     | LI R5, 0         | 0  | 0  |    |   |
| 18           |       |     | LI R8, 0         | 0  | 0  |    |   |
| 19           |       |     | SUB R9, R9, 63   | 0  | 0  |    |   |

As well as in Loop2 and Loop3, all the branches from these loops are backward branches, we can be sure that they are also note with bit number 1 (taken) because after the branches, instructions will instantly repeat from top of Loop2 and Loop3 rather than NOP of instruction following these branches first.

```
.Loop2
    li R4, 0
    li R5, 0
    li R8, 0

    sub R9, R9, 63

    add R11, R11, 1
    sub R29, R11, RLEN
    brnz R29, Loop

    add R7, R7, RLEN
```

```
.Loop3
    li R11, 0
    li R9, SIZE

    add R12, R12, 1
    sub R28, R12, RLEN

    brnz R28, Loop
```

| Memory views |       |     |                    |    |    |    |   |
|--------------|-------|-----|--------------------|----|----|----|---|
| Address      | Label | ... | Instruction        | LI | EX | GP | P |
| 18           |       |     | LI R8, 0           | 1  | 1  |    |   |
| 19           |       |     | SUB R9, R9, 63     | 1  | 1  |    |   |
| 20           |       |     | ADD R11, R11, 1    | 1  | 1  |    |   |
| 21           |       |     | SUB R29, R11, RLEN | 1  | 1  |    |   |
| 22           |       |     | BRNZ R29, LOOP     | 1  | 1  | 1  | T |
| 23           |       |     | ADD R7, R7, RLEN   | 0  | 0  |    |   |
| 24           |       |     | BR LOOP3           | 0  | 0  |    |   |
| 25           | LOOP3 |     | LI R11, 0          | 0  | 0  |    |   |
| 26           |       |     | LI R9, SIZE        | 0  | 0  |    |   |
| 27           |       |     | ADD R12, R12, 1    | 0  | 0  |    |   |
| 28           |       |     | SUB R28, R12, RLEN | 0  | 0  |    |   |

Branch in Loop2

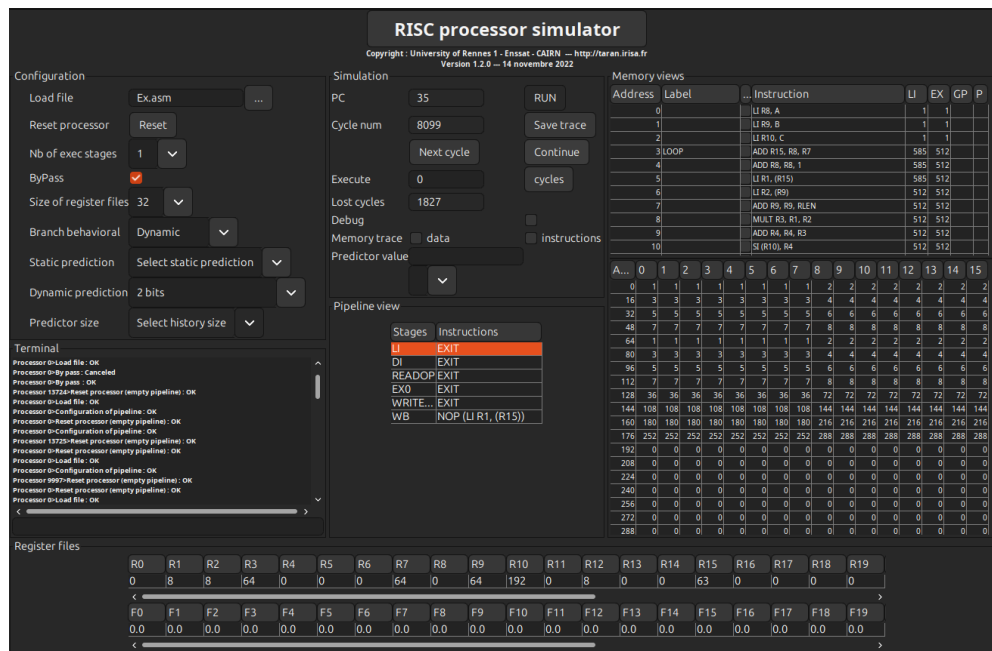
| Memory views |       |     |                    |    |    |    |   |
|--------------|-------|-----|--------------------|----|----|----|---|
| Address      | Label | ... | Instruction        | LI | EX | GP | P |
| 24           |       |     | BR LOOP3           | 1  | 1  |    |   |
| 25           | LOOP3 |     | LI R11, 0          | 1  | 1  |    |   |
| 26           |       |     | LI R9, SIZE        | 1  | 1  |    |   |
| 27           |       |     | ADD R12, R12, 1    | 1  | 1  |    |   |
| 28           |       |     | SUB R28, R12, RLEN | 1  | 1  |    |   |
| 29           |       |     | BRNZ R28, LOOP     | 1  | 1  | 1  | T |
| 30           | EXIT  |     | EXIT               | 0  | 0  |    |   |
| 31           |       |     | EXIT               | 0  | 0  |    |   |
| 32           |       |     | EXIT               | 0  | 0  |    |   |
| 33           |       |     | EXIT               | 0  | 0  |    |   |
| 34           |       |     | EXIT               | 0  | 0  |    |   |

Branch in Loop3

→ The 1-bit predictor provides a simple mechanism for improving the accuracy of branch prediction over static methods. By updating the prediction bit based on the most recent outcome, it can adapt to changing branch behaviors. However, due to its simplicity, it may not handle complex branching patterns efficiently. For example when both backward and forward branches appear in the code multiple times.

## Step 4: Configure dynamic branch prediction, 2 bits (explain how you configure)

- Configure dynamic branch prediction, 2-bit (explain how you configure)



- Analyze the nb of cycles, comments, conclusion

As can be seen from the image, when I enable Dynamic branch behavior technique, the cycle number decreases to 8099 compared with more than 9900 originally. The lost cycles decreased significantly to 1827.

Explanation: A 2-bit saturating counter is used to predict whether a branch will be taken or not taken based on its recent history. Each branch has an associated 2-bit counter that can be in one of four states, representing a strong or weak prediction of either taken or not taken. The counter is updated based on the actual outcome of the branch → 2-bit saturating counters provide a robust and efficient means of dynamic branch prediction. By leveraging recent branch behavior to predict future actions, they significantly improve the accuracy of branch predictions, leading to better utilization of the processor pipeline and overall performance improvements.

- Analyze cycle by cycle

In all my loops, there are only backward branch instructions and no forward branch instructions.

For loop1, loop2 and loop3:

```

Loop
    add R15, R8, R7
    add R8, R8, 1
    li R1, (R15)

    li R2, (R9)
    add R9, R9, RLEN

    mult R3, R1, R2
    add R4, R4, R3
    si (R10), R4

    sub R5, R5, 1
    add R6, R5, RLEN
    brnz R6, Loop

    add R10, R10, 1

```

| Stages   | Instructions        |
|----------|---------------------|
| LI       | ADD R15, R8, R7     |
| DI       | NOP (LI R5, 0)      |
| READOP   | NOP (LI R4, 0)      |
| EX0      | NOP (ADD R10, R1... |
| WRITE... | BRNZ R6, LOOP       |
| WB       | -                   |

**Not using Dynamic  
branch prediction:  
2-bits mode**

| Stages   | Instructions     |
|----------|------------------|
| LI       | LI R2, (R9)      |
| DI       | LI R1, (R15)     |
| READOP   | ADD R8, R8, 1    |
| EX0      | ADD R15, R8, R7  |
| WRITE... | BRNZ R6, LOOP    |
| WB       | ADD R6, R5, RLEN |

**Using Dynamic  
branch prediction:  
2-bits mode**

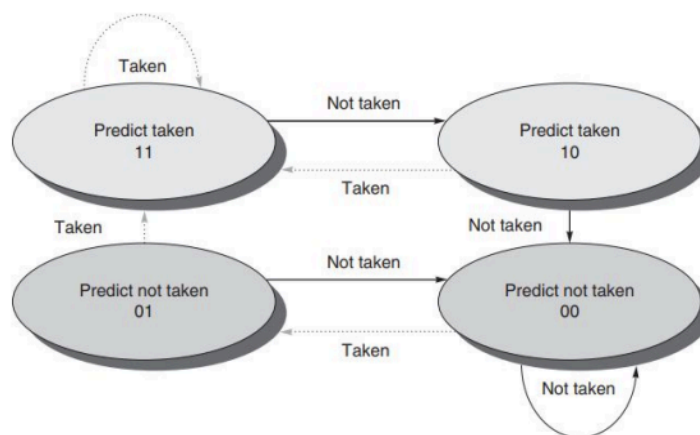
A 2-bit branch predictor employs a two-bit saturating counter to remember the outcome of the last branch instruction. This counter can be in one of four states: "strongly taken," "weakly taken," "weakly not taken," or "strongly not taken." In the given scenario, let's assume the initial state of the counter is "weakly not taken." When the instruction "BRNZ R6, LOOP" is executed, and there are no intervening instructions such as "NOP (ADD R10, R10, 1)" or "NOP (LI R4, 0)" following it before reaching instructions within the loop ("ADD R15, R8, R7"), the counter would transition to a more confident state of prediction, such as "weakly taken" or even "strongly taken" depending on the predictor's history. During subsequent executions, this 2-bit counter would be dynamically updated based on the actual outcomes of branch instructions, thereby refining the predictor's accuracy over time.

| Memory views |       |     |                  |    |    |    |    |
|--------------|-------|-----|------------------|----|----|----|----|
| Address      | Label | ... | Instruction      | LI | EX | GP | P  |
| 9            |       |     | ADD R4, R4, R3   | 16 | 16 |    |    |
| 10           |       |     | SI (R10), R4     | 16 | 16 |    |    |
| 11           |       |     | SUB R5, R5, 1    | 16 | 16 |    |    |
| 12           |       |     | ADD R6, R5, RLEN | 16 | 16 |    |    |
| 13           |       |     | BRNZ R6, LOOP    | 16 | 16 | 14 | FT |
| 14           |       |     | ADD R10, R10, 1  | 2  | 1  |    |    |
| 15           |       |     | BR LOOP2         | 1  | 1  |    |    |
| 16           | LOOP2 |     | LI R4, 0         | 1  | 1  |    |    |
| 17           |       |     | LI R5, 0         | 1  | 1  |    |    |
| 18           |       |     | LI R8, 0         | 1  | 1  |    |    |
| 19           |       |     | SUB R9, R9, 63   | 1  | 1  |    |    |

As well as in Loop2 and Loop3, all the branches from these loops are backward branches, we can be sure that they are also noted with FT because after the branches, instructions will instantly repeat from top of Loop2 and Loop3 rather than NOP of instruction following these branches first.

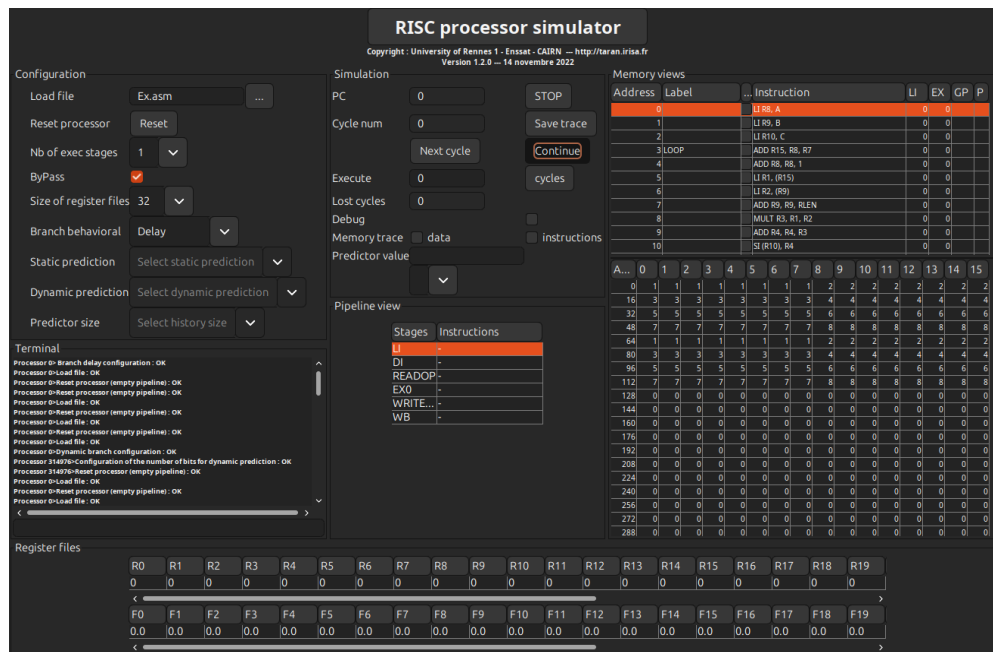
| Memory views |       |                    |    |    |    |    |
|--------------|-------|--------------------|----|----|----|----|
| Address      | Label | ... Instruction    | LI | EX | GP | P  |
| 21           |       | SUB R29, R11, RLEN | 64 | 64 |    |    |
| 22           |       | ✓ BRNZ R29, LOOP   | 64 | 64 | 56 | FT |
| 23           |       | ADD R7, R7, RLEN   | 8  | 8  |    |    |
| 24           |       | BR LOOP3           | 8  | 8  |    |    |
| 25           | LOOP3 | LI R11, 0          | 8  | 8  |    |    |
| 26           |       | LI R9, SIZE        | 8  | 8  |    |    |
| 27           |       | ADD R12, R12, 1    | 8  | 8  |    |    |
| 28           |       | SUB R28, R12, RLEN | 8  | 8  |    |    |
| 29           |       | ✓ BRNZ R28, LOOP   | 8  | 8  | 7  | FT |
| 30           | EXIT  | EXIT               | 1  | 0  |    |    |
| 31           |       | EXIT               | 1  | 0  |    |    |

→ Each branch in the program is associated with a 2-bit saturating counter. This counter can take four states: strongly taken, weakly taken, weakly not taken, and strongly not taken. Initially, all counters are set to a default state, often "weakly not taken" or "weakly taken". It maintains a history of branch behavior using 2-bit saturating counters associated with each branch instruction. This history helps in making accurate predictions about whether a branch will be taken or not taken in the future. Overall, 2-bit dynamic branch prediction significantly enhances the performance of processors by reducing branch misprediction penalties and improving overall execution efficiency.



## Step 5: Analyze branch prediction techniques (delay branch)

### - Configure delay branch



As can be seen from the image, when I choose Delay branch behavior, the execution of the assembly code did not finish. The execution takes forever and there is no result is written back to register or memory.

### - What do you have to do to ensure execution of the code

When implementing the delayed branch technique, the processor executes the conditional branch instruction in the usual manner but postpones the calculation of the branch target address. Instead, it proceeds to execute the subsequent instructions following the branch instruction until the branch target address is determined. Upon computing the branch target address, the processor retrieves the instruction located at that address and initiates its execution.

→ We have to modify the code (adding 3 NOPs following each branch “BRNZ” because there are 3 instructions are executed after them) in order to ensure that when the branch is executed, the following instruction is not executed (or stop executed) until the branch target address is determined.

### - Modify the code, and verify the execution



```

Loop
    add R15, R8, R7    -- R15 = R8 + R7 (compute address offset for current element in a)
    add R8, R8, 1      -- Increment pointer for a
    li R1, (R15)       -- Load current element of a into R1

    li R2, (R9)        -- Load current element of b into R2
    add R9, R9, RLEN    -- Move to the next row in matrix b

    mult R3, R1, R2     -- Multiply R1 and R2, store result in R3
    add R4, R4, R3      -- Add result to R4 (accumulator)
    si (R10), R4        -- Store accumulator result in current position of c

    sub R5, R5, 1       -- Decrement R5 (loop counter)
    add R6, R5, RLEN     -- R6 = R5 + RLEN
    brnz R6, Loop       -- If R6 is not zero, continue Loop
    nop
    nop
    nop

    add R10, R10, 1     -- Increment pointer for c
    br Loop2

Loop2
    li R4, 0            -- Set R4 to 0
    li R5, 0            -- Set R5 to 0
    li R8, 0            -- Set R8 to 0

    sub R9, R9, 63      -- R9 = R9 - (SIZE - 1) to move to next column of b

    add R11, R11, 1     -- Increment column counter
    sub R29, R11, RLEN  -- R29 = R11 - RLEN
    brnz R29, Loop     -- If R29 is not zero, continue Loop
    nop
    nop
    nop

    add R7, R7, RLEN    -- Move to next row of a
    br Loop3

Loop3
    li R11, 0           -- Reset column counter
    li R9, SIZE         -- Reset pointer for b to start of the next column

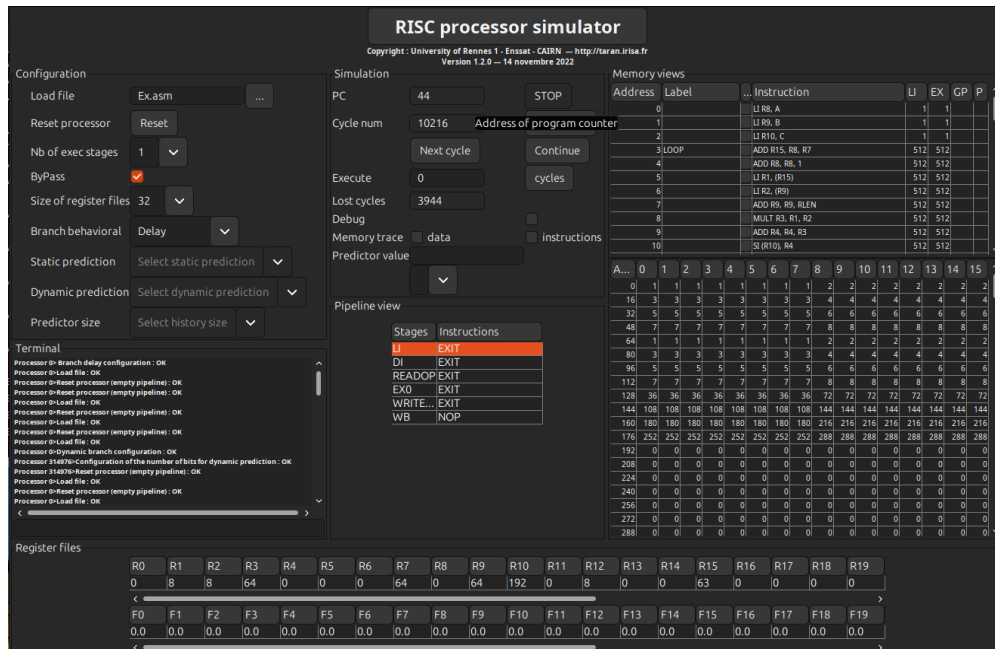
    add R12, R12, 1     -- Increment row counter
    sub R28, R12, RLEN  -- R28 = R12 - RLEN

    brnz R28, Loop     -- If R28 is not zero, continue Loop
    nop
    nop
    nop

```

As can be seen from the image above, 3 NOP instructions have been added after each BRNZ instruction in order to skip the execution of 3 instructions after these branches.

Execution after modifying the code:



## - Analyze cycle by cycle

When implementing the delayed branch technique, the processor executes the conditional branch instruction in the usual manner but postpones the calculation of the branch target address. Instead, it proceeds to execute the subsequent instructions following the branch instruction until the branch target address is determined. Upon computing the branch target address, the processor retrieves the instruction located at that address and initiates its execution.

Through overlapping the execution of instructions that come after the branch instruction with the computation of the branch target address, the delayed branch technique aims to minimize the number of cycles lost due to branch mispredictions. Nevertheless, this approach can lead to an increase in the overall program execution time, as it introduces additional processing overhead for computing the branch target address and managing the coordination between instructions.

| Stages   | Instructions    |
|----------|-----------------|
| LI       | ADD R15, R8, R7 |
| DI       | NOP             |
| READOP   | NOP             |
| EXO      | NOP             |
| WRITE... | BRNZ R6, LOOP   |
| WB       | -               |

| Stages   | Instructions    |
|----------|-----------------|
| LI       | ADD R15, R8, R7 |
| DI       | LI R4, 0        |
| READOP   | BR LOOP2        |
| EXO      | ADD R10, R10, 1 |
| WRITE... | BRNZ R6, LOOP   |
| WB       | -               |

In the scenario depicted in the images, when I introduced three NOP instructions after the "BRNZ R6, LOOP" instruction, the behavior changed. In the original code, following the execution of "BRNZ R6, LOOP", instructions beyond the branch, such as "ADD R10, R10, 1" and "BR Loop2", were executed. This led to an undesired outcome and an infinite loop due to the premature execution of instructions from the Loop1 block.

However, when the code was modified to include NOP instructions after the branch, the processor engaged in speculative execution, continuing with instructions beyond the branch while simultaneously initiating execution from the target of the branch. This speculative execution prevented the premature execution of instructions from the Loop1 block, potentially avoiding the infinite loop scenario observed in the original code. This behavior exemplifies how delayed branch execution can assist in maintaining program correctness and avoiding unintended outcomes by allowing the processor to defer the resolution of branches until after speculative execution has progressed.

## - Modify the nb of execution stages

After I change the nb of execution stages to 3, the result is wrong again:

The screenshot displays the RISC processor simulator interface. The configuration panel on the left shows the number of execution stages set to 3. The simulation panel in the center shows the PC at 8 and cycle number at 2232. The memory views panel on the right shows the instruction stream, with the current instruction being "ADD R9, R9, RLEN" at address 7. The register files at the bottom show the state of the processor registers.

**RISC processor simulator**  
Copyright : University of Rennes 1 - Enstat - CARM - http://taran.irisa.fr  
Version 1.2.0 - 14 novembre 2022

**Configuration**

- Load file: Exasm
- Reset processor: Reset
- Nb of exec stages: 3
- ByPass: ☒ Number of stages for the execution units: 0
- Size of register files: 32
- Branch behavior: Delay
- Static prediction: Select static prediction
- Dynamic prediction: Select dynamic prediction
- Predictor size: Select history size

**Simulation**

- PC: 8
- Cycle num: 2232
- Next cycle
- Lost cycles: 1312
- Debug: ☐ data, ☐ instructions
- Memory trace: ☐ data, ☐ instructions
- Predictor value:

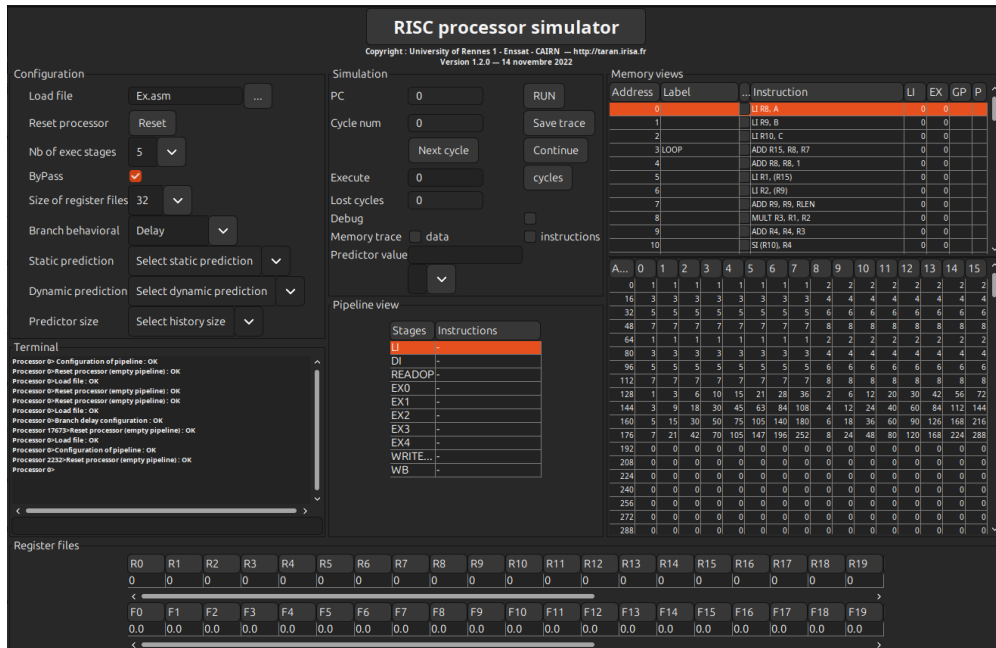
**Memory views**

| Address | Label | Instruction      | LI | EX | GP | P  |
|---------|-------|------------------|----|----|----|----|
| 0       |       | LI R8, A         | 1  | 1  |    |    |
| 1       |       | LI R9, B         | 1  | 1  |    |    |
| 2       |       | LI R10, C        | 1  | 1  |    |    |
| 3       | LOOP  | ADD R15, R8, R7  |    |    | 65 | 64 |
| 4       |       | ADD R8, R8, 1    |    |    | 65 | 64 |
| 5       |       | LI R1, (R15)     |    |    | 65 | 64 |
| 6       |       | LI R2, (R9)      |    |    | 65 | 64 |
| 7       |       | ADD R9, R9, RLEN |    |    | 65 | 64 |
| 8       |       | MULT R3, R1, R2  |    |    | 64 | 64 |
| 9       |       | ADD R4, R4, R3   |    |    | 64 | 64 |
| 10      |       | LI R10, R4       |    |    | 64 | 64 |

**Register files**

| R0 | R1 | R2 | R3 | R4 | R5 | R6 | R7 | R8 | R9 | R10 | R11 | R12 | R13 | R14 | R15 | R16 | R17 | R18 | R19 |
|----|----|----|----|----|----|----|----|----|----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0  | 8  | 8  | 64 | 0  | 0  | 0  | 64 | 0  | 64 | 192 | 0   | 1   | 0   | 0   | 63  | 0   | 0   | 0   | 0   |

When I change the nb of execution stages to 5, the execution is not work due to infinite loop:



## Step 6: Analyze branch prediction techniques

### - Conditional move

Conditional move (CMOV) is a branch prediction technique that aims to mitigate the performance impact of conditional branches by replacing them with conditional move instructions where possible. Rather than branching based on a condition and executing different instructions depending on the outcome, conditional move instructions conditionally move data into a destination register based on a condition, typically without altering the control flow.

Syntax of Conditional move:

|                             |  |
|-----------------------------|--|
| cmovz, cmovnz, cmovn, cmovp | Conditionnal copy from one register to another registres |
|-----------------------------|--|

The objective of this section is to attempt to replace certain conditional branch instructions with conditional moves. But JSimRISC Processor Simulation does not have conditional move instructions (means that syntax in the image above is not working in JSimRisc Processor Simulator), and even if it did, it is impossible to replace all conditional branch instructions by conditional move instructions because conditional moves can only work for complex branching logic.